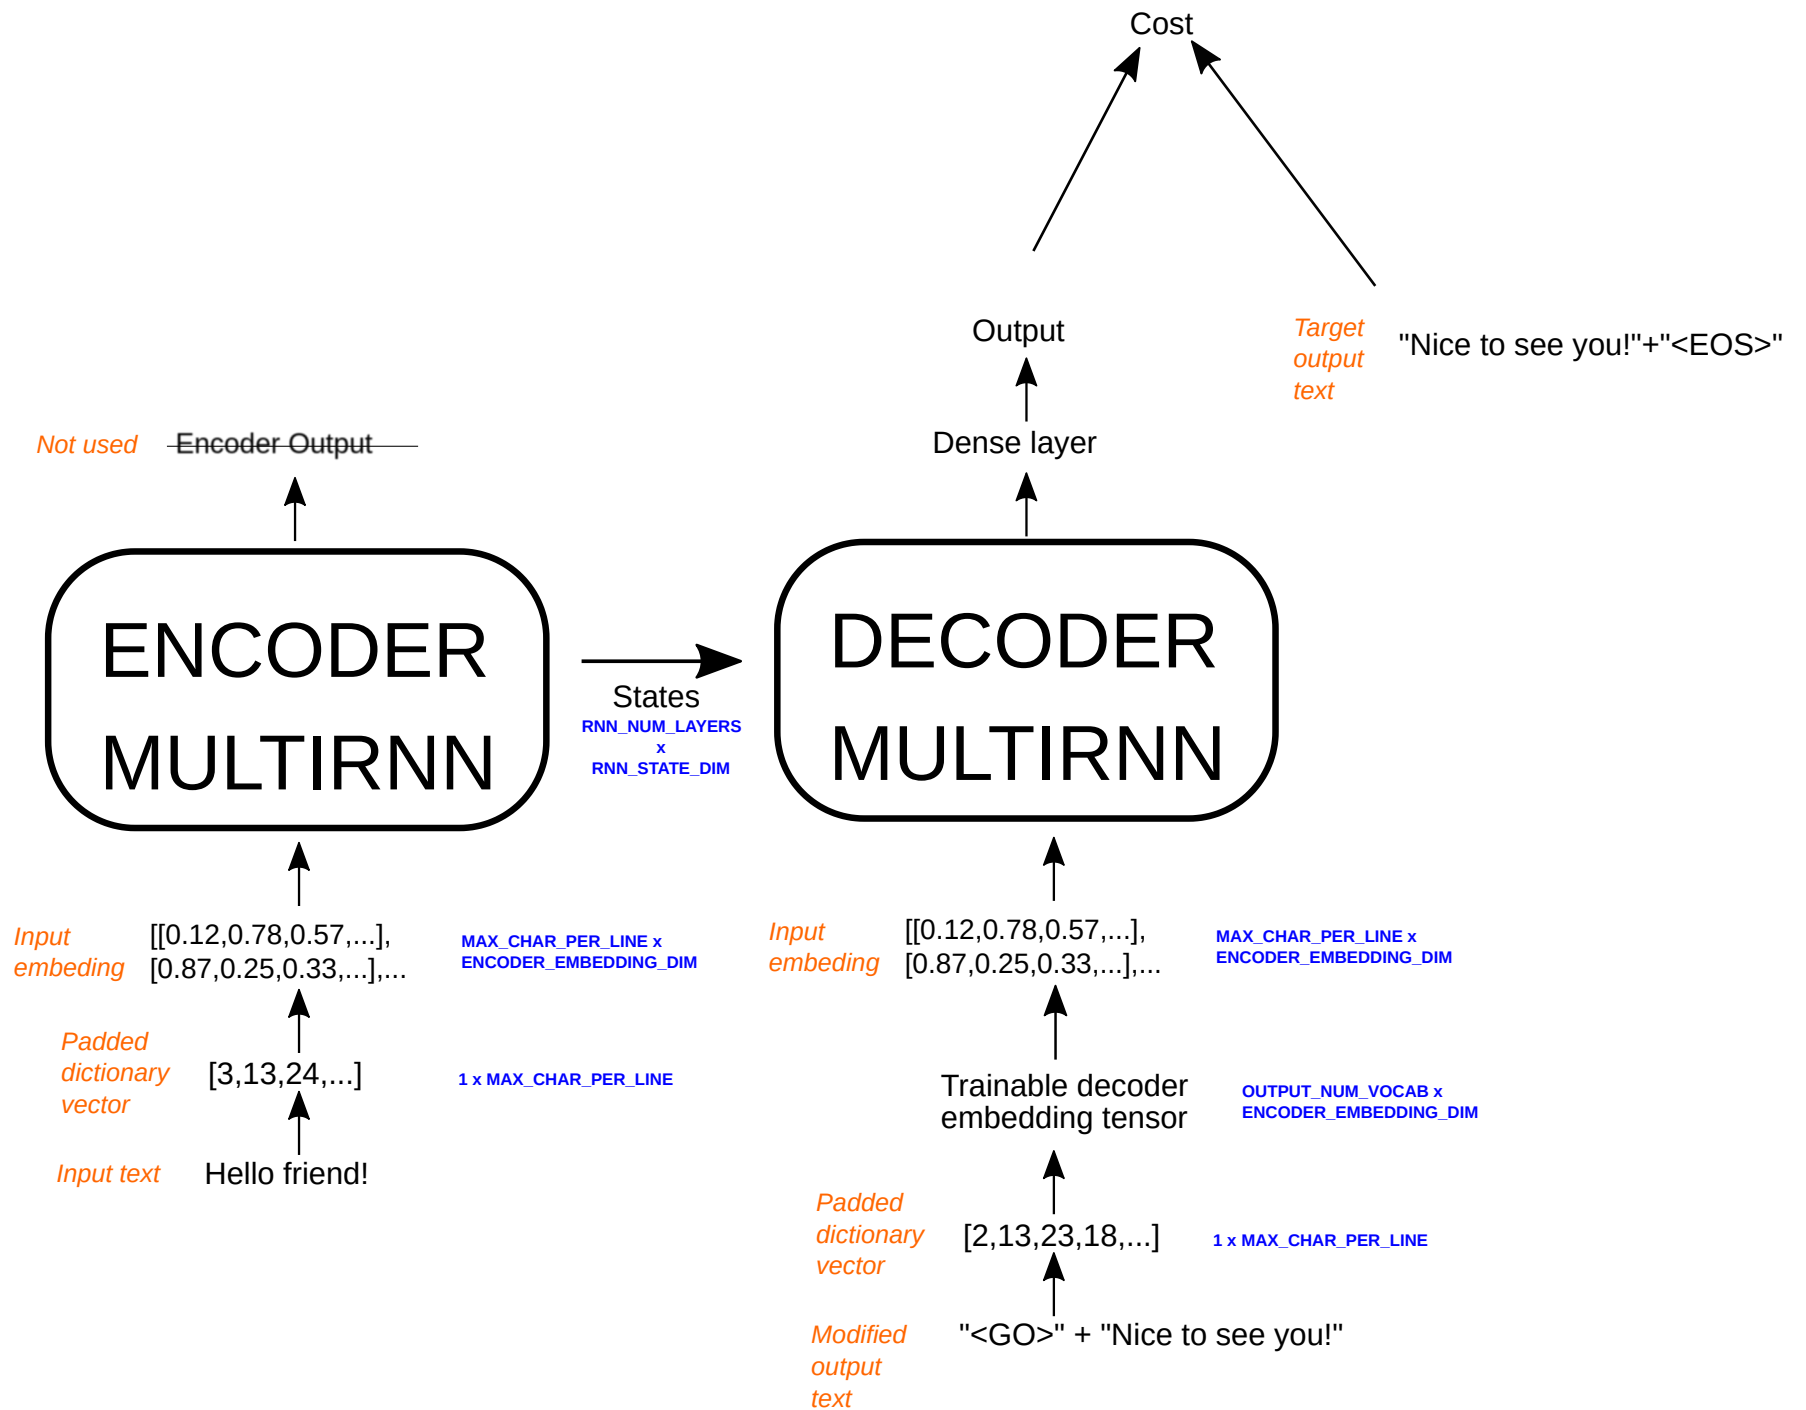


Step by step example

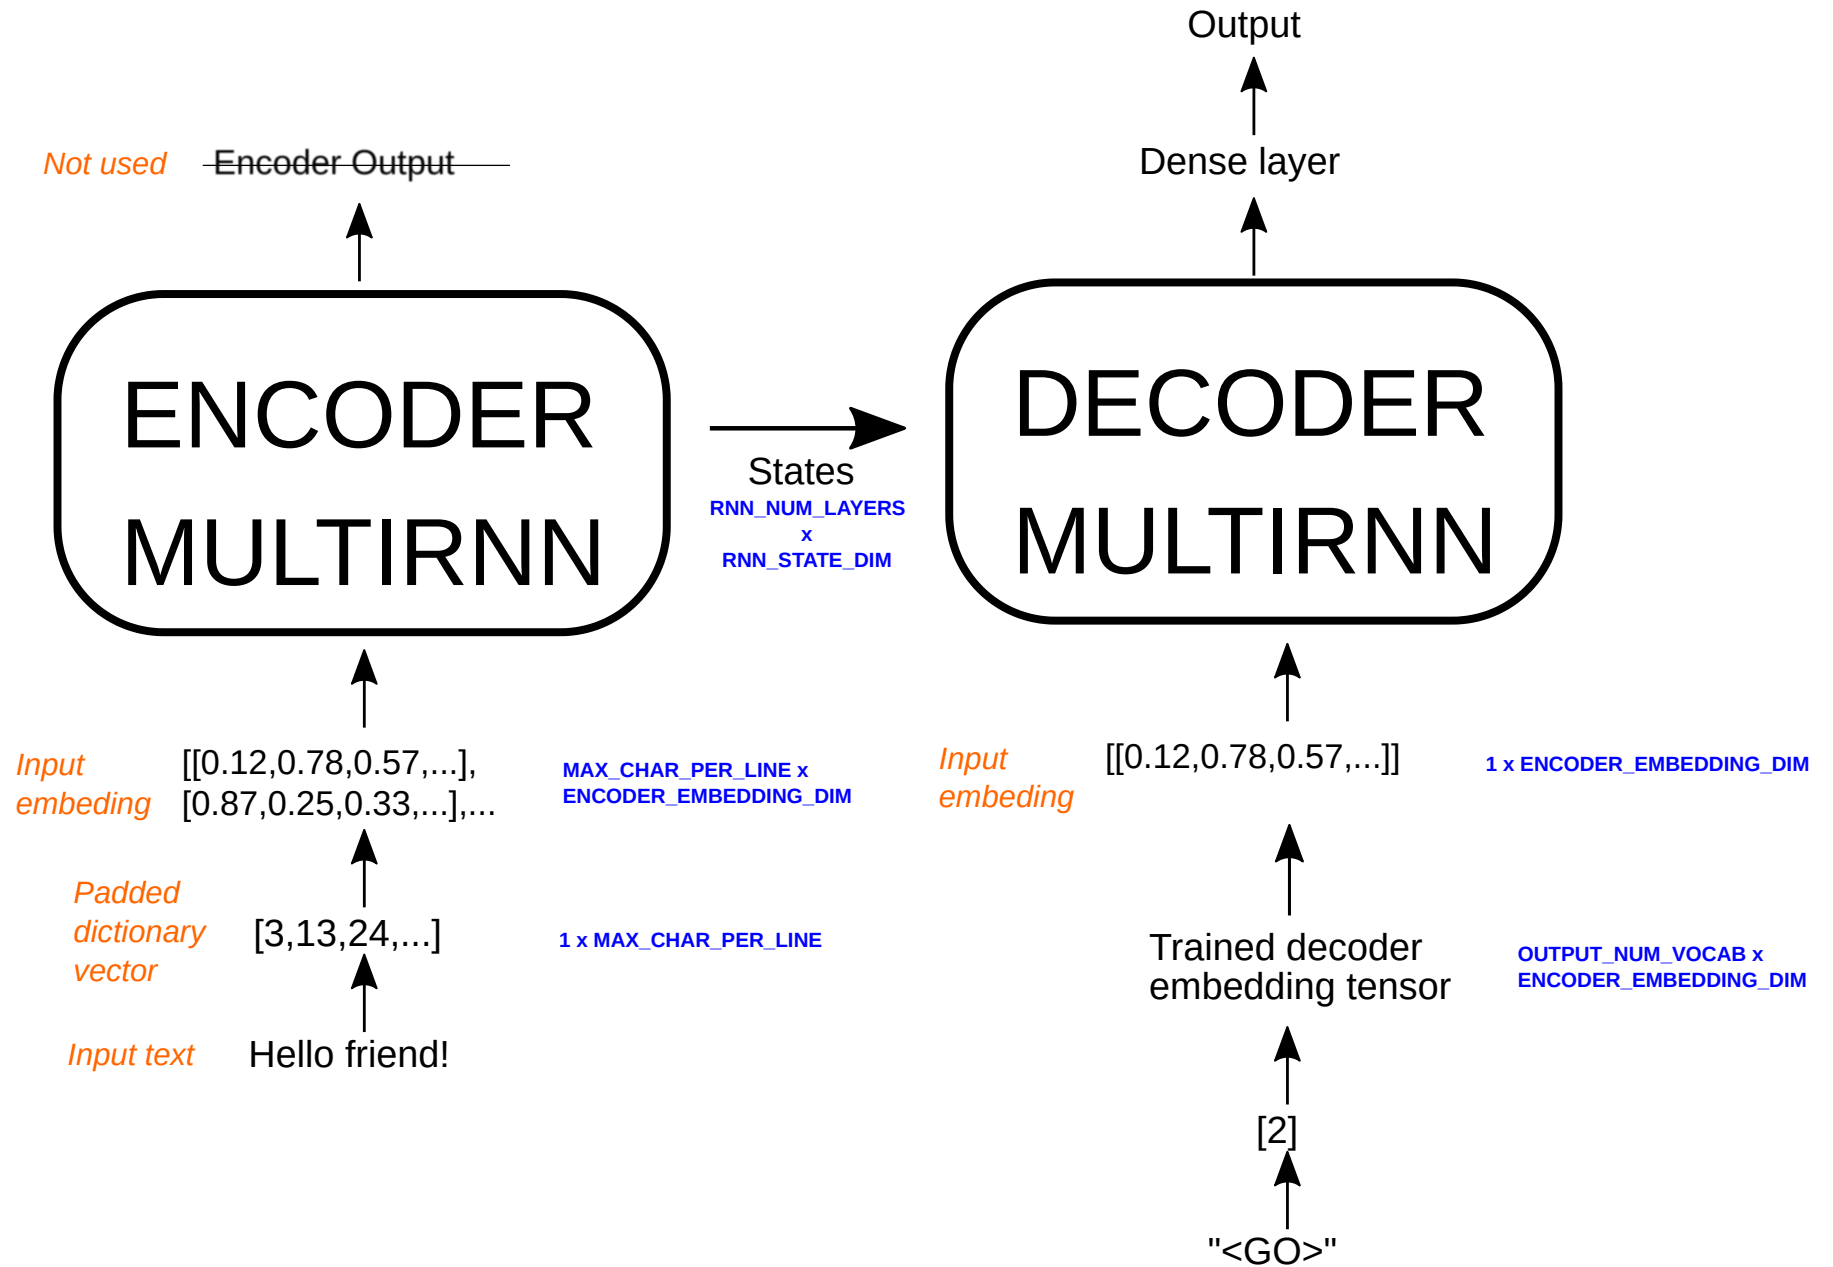
Practical explanation

Overview

Training



Inference



Inputs/Outputs

When we train a seq2seq model we pass input and output sentences. These sentences will need to be cropped to a maximum length limit (`MAX_CHAR_PER_LINE`) and also each character assigned to a dictionary index.

This is achieved by the functions `load_sentences` and `extract_character_vocab` .

For example if we have sentence 'some sentence' it will become a vector representation with the dictionary indices `[2, 3, 4, 5, 6, ...]` , this vector will have a maximum size of `MAX_CHAR_PER_LINE` .

If sentences are shorter than `MAX_CHAR_PER_LINE` the remainder of the vector representation is filled with the dictionary index for the character `<PAD>` , using the function `pad` .

In the code example below the two sentences are

```
Input sentence = 'she okay'
vector: [20, 30, 22, 5, 23, 12, 7, 19, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]

Output sentence = 'i hope so'
vector: [13, 5, 30, 23, 27, 22, 5, 20, 23, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

Encoder input embeddings

The input at this point is a sequence of integers. Our model will benefit from a denser representation of the input using embeddings compared to one-hot representation of each of the integer.

In the code this embedding is achieved by

```
encoder_input_embedded = tf.contrib.layers.embed_sequence(
    ids=padded_symbols_input, # current input sequence of numbers
    vocab_size=INPUT_NUM_VOCAB, # rows of embedding matrix from dictionary size
    embed_dim=ENCODER_EMBEDDING_DIM # cols of embedding matrix from user
)
```

For the input in the example above it becomes:

```
[[ 0.15645272  0.20985416 -0.10991608 ...  0.09195071  0.05068645 -0.15631753]
 [-0.05581941  0.16376457 -0.12799817 ... -0.14128818 -0.07264952 -0.17910267]
 [-0.1667774  -0.04544044  0.13017449 ...  0.06774965  0.02339229 -0.10449551]
 ...
 [ 0.16037011  0.19128785  0.05270347 ...  0.23134735  0.1546481 -0.18897504]
 [ 0.16037011  0.19128785  0.05270347 ...  0.23134735  0.1546481 -0.18897504]
 [ 0.16037011  0.19128785  0.05270347 ...  0.23134735  0.1546481 -0.18897504]]
```

where each row is a vector of size `ENCODER_EMBEDDING_DIM` representing each integer in

```
[14, 29, 25, 13, 6, 30, 17, 12, 13, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

See <https://towardsdatascience.com/neural-network-embeddings-explained-4d028e6f0526> (<https://towardsdatascience.com/neural-network-embeddings-explained-4d028e6f0526>) for more

Encoder

The encoder is a multiRNN as explained in `Chapter_11_Section_1_MultiRNN.ipynb`. The functions used to create this multiRNN are `make_cell`, `make_multi_cell`

The output of the encoder is obtained by running the embedded input through `dynamic_rnn` tensorflow function:

```
encoder_output, encoder_state = tf.nn.dynamic_rnn(encoder_multi_cell,
                                                  encoder_input_embedded,
                                                  sequence_length=(len(padded_symbols_input),),
                                                  dtype=tf.float32)
```

For more see `Chapter_11_Section_1_MultiRNN.ipynb`

The encoder output is deleted (`del encoder_output`) and the state is passed into the decoder.

The encoder state has dimensions `RNN_NUM_LAYERS × RNN_STATE_DIM` and is printed out in the code example with

```
for i, encoder_state in enumerate(encoder_state):
    print('Layer {}, hidden state shape {}'.format(i, encoder_state[0].shape))
    print('Layer {}, activation state shape {}'.format(i, encoder_state[1].shape))
```

Decoder

Training input

The training input to the decoder is a modified version of the output

```
output_sentence = 'i hope so'+padding
vector: [13, 5, 30, 23, 27, 22, 5, 20, 23, 5, 0, 0, 0, 0, 0, 0, 0, 0]
```

Where we add the `'<G0>'` vocabulary representation at the beggining and clip to the `MAX_CHAR_PER_LINE`

```
Input sentence = '<G0>' + clipped_output_sentence
vector: [2, 13, 5, 30, 23, 27, 22, 5, 20, 23, 5, 0, 0, 0, 0, 0, 0, 0, 0]
```

Training output

For training notice that we attach a <EOS> at the end of each sentence.

```
# output_sentence = 'i hope so'+ '<EOS>' + padding
symbols_output[-1] = input_symbol_to_int['<EOS>']
```

Inference input

For training it's ok to use a modified version of the output to assist with training (since we have the ground truth) but for inference (final predictions) the input to the decoder will just be the vocabulary representation of <G0> . The code to produce the inference inputs is:

```
start_tokens = tf.tile(tf.constant([output_symbol_to_int['<G0>']],
                                   dtype=tf.int32), [BATCH_SIZE], name='start_tokens')
```

Decoder embedding

The decoder embedding is a trainable tensor. The aim is to learn good embedding representations that based on the encoder state, and the inference input <G0> it will produce reasonable outputs for a conversation.

The first dimension of the `decoder_embedding` variable tensor represents the number of entries in our output vocabulary and it's size is `OUTPUT_NUM_VOCAB` , the second dimension is user defined and in the example code is defined in variable `DECODER_EMBEDDING_DIM` . The definition of the encoder embedding in the code is initialized by:

```
decoder_embedding = tf.Variable(tf.random_uniform([OUTPUT_NUM_VOCAB, DECODER_EMBEDDING_DIM]))
```

For a specific input to the decoder, for example for training this would be the vocabulary sequence of '<G0>' + `clipped_output_sentence` , we can extract it's embedding using:

```
decoder_input_embedded = tf.nn.embedding_lookup(decoder_embedding, decoder_input_seq)
```

where `decoder_input_seq` is the vocabulary sequence of '<G0>' + `clipped_output_sentence`

Training Decoder MultiRNN

The decoder is constructed in a more complicated way than the encoder multiRNN. Instead of using `tf.nn.dynamic_rnn` we use `tf.contrib.seq2seq.BasicDecoder`. This is done because in the case of the decoder we need to handle the input states from the encoder, we also need to use a fully connected `Dense` layer to transform the output of the decoder multiRNN to a one-hot representation of the output vocabulary and also `tf.contrib.seq2seq.BasicDecoder` accepts a `tf.contrib.seq2seq.TrainingHelper` which will handle the input to the decoder for us.

multiRNN cell

Same as with the encoder

```
decoder_multi_cell = make_multi_cell(RNN_STATE_DIM, RNN_NUM_LAYERS)
```

Fully connected Dense layer

```
output_layer = Dense(OUTPUT_NUM_VOCAB, kernel_initializer=tf.truncated_normal_initializer(mean=0.0, stddev=0.1))
```

Input training helper

The `tf.contrib.seq2seq.TrainingHelper` will handle the input to the decoder for us

```
training_helper = tf.contrib.seq2seq.TrainingHelper(inputs=decoder_input_embedded,  
                                                    sequence_length=(len(padded_symbols_output)), time_major=False)
```

Decoder main

The `tf.contrib.seq2seq.BasicDecoder` is essentially what `tf.nn.dynamic_rnn` was for the encoder. It accepted the multiRNN cells, the `training_helper` that handles the input, the encoder state and the `Dense` `output_layer`.

```
training_decoder = tf.contrib.seq2seq.BasicDecoder(decoder_multi_cell,  
                                                    training_helper, encoder_state, output_layer)
```

Decoder outputs

The `tf.contrib.seq2seq.dynamic_decode` produces the outputs of the decoder as `[final_outputs, final_state, final_sequence_lengths]`

```
training_decoder_output_seq, _, _ = tf.contrib.seq2seq.dynamic_decode(training_decoder,  
                                                                        impute_finished=True, maximum_iterations=len(padded_symbols_output))
```

We can access the output from the decoder using `training_decoder_output_seq.rnn_output` which is a tensor with dimensions `[batchSize, length_current_output, size_output_dictionary]`. `length_current_output` is defined by `MAX_CHAR_PER_LINE`.

So for each output element we produce a vector of length `size_output_dictionary` (`[:,i,:]`), from that we select the highest (?) activation to be the vocabulary character to output.

Inference Decoder MultiRNN

The inference decoder is build in a similar way as the training decoder.

```
# Helper for the inference process. It takes the whole of decoder embeddings
# and the integer representation of the <EOS> symbol, which is the 'end_token'
# This helper will handle the inputs to the decoder for inference.
inference_helper = tf.contrib.seq2seq.GreedyEmbeddingHelper(embedding=decoder_embedding,
    start_tokens=start_tokens,end_token=output_symbol_to_int['<EOS>'])
# Basic decoder for inference. It uses the decoder MultiRNN 'decoder_multi_cell'
# inference_helper takes care of the input (input only the <GO> symbol), the encoder state
# which takes the sentence for which we want to have a reply from our chatbot
# and the output layer that turns the output into a one-hot representation corresponding to
# our vocabulary
inference_decoder = tf.contrib.seq2seq.BasicDecoder(decoder_multi_cell,
    inference_helper,encoder_state,output_layer)
# Perform dynamic decoding using the decoder
# this is the output, same as in the case of the train decoder
inference_decoder_output_seq, _, _ = tf.contrib.seq2seq.dynamic_decode(inference_decoder,
    impute_finished=True,maximum_iterations=MAX_CHAR_PER_LINE)
```

Cost

The ground truth is compared to the output of the model. For this we are going to use `tf.contrib.seq2seq.sequence_loss` which takes the output of the training decoder `training_decoder_output_seq.rnn_output` , the tensor with the ground truth and a mask that controls which parts of the output/ground truth sequences to compare. For example we might want to use the mask hide the `<PAD>` symbol so the model doesn't learn to pad a sentence.

The code to calculate cost of the example is:

```
outputGroundTruth=tf.convert_to_tensor(padded_symbols_output,dtype=tf.int32)
# add batch dummy index
batchGroundTruth=tf.expand_dims(outputGroundTruth,axis=0)
# mask is used to clip the output sequence to its maximum allows size
masks = tf.sequence_mask(tf.constant([len(symbols_output)]),MAX_CHAR_PER_LINE,
    dtype=tf.float32,name='masks')
cost = tf.contrib.seq2seq.sequence_loss(training_decoder_output_seq.rnn_output,
    batchGroundTruth,masks)
```


Code

Import libraries

```
In [1]: 1 import tensorflow as tf
2 import numpy as np
3 import os, re
4 from tensorflow.python.layers.core import Dense
```

In the example code we are just going to run through a simple input/output example and calculate the output. So first load the data and create a simple example

```
In [2]: 1 # Functions to load data and create vocabularies
2 def load_sentences(path):
3     with open(path, 'r', encoding="ISO-8859-1") as f:
4         data_raw = f.read().encode('ascii', 'ignore').decode('UTF-8').lower()
5         data_alpha = re.sub('[^a-z\n]+', ' ', data_raw)
6         data = []
7         for line in data_alpha.split('\n'):
8             data.append(line[:MAX_CHAR_PER_LINE])
9     return data
10
11 def extract_character_vocab(data):
12     special_symbols = ['<PAD>', '<UNK>', '<GO>', '<EOS>']
13     # extract unique characters from all sentences in data
14     set_symbols = set([character for line in data for character in line])
15     # add special symbols
16     all_symbols = special_symbols + list(set_symbols)
17     # create vocabularies that match symbol to index, and index to symbol
18     # breakdown of syntax -> dict = {key:item for key, item in enumerate(list)}
19     int_to_symbol = {word_i: word for word_i, word in enumerate(all_symbols)}
20     symbol_to_int = {word: word_i for word_i, word in int_to_symbol.items()}
21     return int_to_symbol, symbol_to_int
22
23 def pad(xs, size, pad):
24     return xs + [pad] * (size - len(xs))
```

Load data and create vocabularies

```
In [3]: 1 MAX_CHAR_PER_LINE = 20
2 # read text as two lists `input_sentences` and `output_sentences`
3 # the two lists have the same length and they have 1:1 correspondence
4 # i.e. `input_sentences[0]`->`output_sentences[0]`, `input_sentences[1]`->`output_sentences[1]`
5 input_sentences = load_sentences(
6     './data/words_input.txt')
7 output_sentences = load_sentences(
8     './data/words_output.txt')
9
10 input_int_to_symbol, input_symbol_to_int = extract_character_vocab(input_sentences)
11 output_int_to_symbol, output_symbol_to_int = extract_character_vocab(output_sentences)
```

Pick two sentences to create an example of input/output

```
In [4]: 1 # input sentece
2 # for i in range(100):
3 #     print('{} , In: {}, out: {}'.format(i, input_sentences[i],output_sentences[i]))
4 exampleIndex=42
5 symbols_input = [input_symbol_to_int[symbol] for symbol in input_sentences[exampleIndex]]
6 padded_symbols_input = pad(symbols_input, MAX_CHAR_PER_LINE, input_symbol_to_int['<PAD>'])
7 print('Input sentence: {}, vocabulary representation: {}'.format(input_sentences[exampleIndex]
8     , padded_symbols_input))
9 symbols_output = [input_symbol_to_int[symbol] for symbol in output_sentences[exampleIndex]]
10 # for training an <EOS> is attached to the end of the output
11 symbols_output[-1]=input_symbol_to_int['<EOS>']
12 padded_symbols_output = pad(symbols_output, MAX_CHAR_PER_LINE, input_symbol_to_int['<PAD>'])
13 print('Output sentence vocabulary representation without padding: {}'.format(symbols_output))
14 print('Output sentence: {}, vocabulary representation: {}'.format(output_sentences[exampleIndex]
15     , padded_symbols_output))
16
```

Input sentence: she okay , vocabulary representation: [23, 13, 19, 28, 25, 4, 11, 30, 28, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]

Output sentence vocabulary representation without padding: [14, 28, 13, 25, 24, 19, 28, 23, 25, 3]

Output sentence: i hope so , vocabulary representation: [14, 28, 13, 25, 24, 19, 28, 23, 25, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]

Some constants to be used by our seq2seq model

```
In [5]: 1 # number of neurons of LSTM cell
2 # if we feed a sequence of [batch_size, sequence_length, features] -> [1, 3, 1]
3 # we will get back [batch_size, sequence_length, RNN_STATE_DIM] -> [1, 3, 512]
4 # of course if we are interested only in the last output we can get it with[:, -1, :]
5 RNN_STATE_DIM = 512
6 # number of layers of the multiRNN cell, see `Concept01_multi_rnn_DC.ipynb`
7 RNN_NUM_LAYERS = 2
8 # number of dimensions of embedding layer
9 ENCODER_EMBEDDING_DIM = DECODER_EMBEDDING_DIM = 64
10 # sizes of input/output vocabularies
11 INPUT_NUM_VOCAB = len(input_symbol_to_int)
12 OUTPUT_NUM_VOCAB = len(output_symbol_to_int)
```

Helper functions to create MultiRNNCell, see Concept01_multi_rnn_DC.ipynb

```
In [6]: 1 # create LSTM cell
2 def make_cell(state_dim):
3     lstm_initializer = tf.random_uniform_initializer(-0.1, 0.1)
4     return tf.contrib.rnn.LSTMCell(state_dim, initializer=lstm_initializer)
5
6 # create many LSTM cells
7 def make_multi_cell(state_dim, num_layers):
8     cells = [make_cell(state_dim) for _ in range(num_layers)]
9     return tf.contrib.rnn.MultiRNNCell(cells)
```

Model

```

In [7]: 1 config = tf.ConfigProto(allow_soft_placement = True)
2 sess = tf.Session(config = config)
3 with sess:
4     # initialise
5     sess.run(tf.global_variables_initializer())
6     sess.run(tf.local_variables_initializer())
7     # Create embeddings from input of sequence of integers
8     encoder_input_embedded = tf.contrib.layers.embed_sequence(
9         ids=padded_symbols_input, # input seq of numbers (row indices)
10        vocab_size=INPUT_NUM_VOCAB, # rows of embedding matrix
11        embed_dim=ENCODER_EMBEDDING_DIM # cols of embedding matrix
12    )
13    sess.run(tf.global_variables_initializer())
14    sess.run(tf.local_variables_initializer())
15    # **** INPUT EMBEDDING ****
16    print('*****Input*****')
17    print('Length of input sequence: {}'.format(len(padded_symbols_input)))
18    print('*****Embedded Input*****')
19    print('Shape of embeded input -> rows=lengthInputSequence={}, \
20        columns=userDefined=ENCODER_EMBEDDING_DIM`={}'.format(encoder_input_embedded.shape[0],
21        encoder_input_embedded.shape[1]))
22    # add degenerate dimension to emulate batch size
23    encoder_input_embedded = tf.expand_dims(encoder_input_embedded, axis=0)
24    print('Reshaped embedded input -> [batch_size, length_of_sequence, \
25        embedding_dimensions]=[{}, {}, {}]'.format(encoder_input_embedded.shape[0],
26        encoder_input_embedded.shape[1], encoder_input_embedded.shape[2]))
27    # **** ENCODER ****
28    encoder_multi_cell = make_multi_cell(RNN_STATE_DIM, RNN_NUM_LAYERS)
29    encoder_output, encoder_state = tf.nn.dynamic_rnn(encoder_multi_cell,
30        encoder_input_embedded, sequence_length=(len(padded_symbols_input)),
31        dtype=tf.float32)
32    print('*****Output-States of Encoder*****')
33    for i, temp_encoder_state in enumerate(encoder_state):
34        print('Layer {}, hidden state shape {}'.format(i, temp_encoder_state[0].shape))
35        print('Layer {}, activation state shape {}'.format(i, temp_encoder_state[1].shape))
36    # we don't need the encoder output only the state
37    del encoder_output
38    # **** DECODER ****
39    # For the decoder we need to pass:
40    # 1) the states from the encoder
41    # 2) the input will be of two kinds: a) one for training that will be the same as the output
42    # sequence except the last sequence item will be removed the first sequence item will be a <G0> symbol
43    # in front so that the decoder will start producing the output without knowing the output
44    # b) the other kind of input is for the inference stage, this will not have any information
45    # about the output so it will only be the <G0> symbol
46    # remove last symbol from input (which for training is a modified output)

```

```

47 decoder_raw_seq = padded_symbols_output[:-1]
48 decoder_input_seq = [output_symbol_to_int['<GO>']] + decoder_raw_seq
49 print('Unmodified output to decoder (training): {}'.format(padded_symbols_output))
50 print('Modified output for input to decoder (training): {}'.format(decoder_input_seq))
51 print('*****Trainable decoder embeddings*****')
52 # initialize embedding vector representations of the input to the decoder
53 # this is initialized to random numbers and will be trained by the seq2seq model
54 # it represents every symbol in out vocabulary (row) and it's vector representation (column)
55 decoder_embedding = tf.Variable(tf.random_uniform([OUTPUT_NUM_VOCAB, # number of rows -> total number of output sy
56                                                    # in output dictionary
57                                                    DECODER_EMBEDDING_DIM # the length of the embedding vectors (hyp
58                                                    ]))
59 print('Shape of decoder embeddings are: row=sizeOutputVocabulary={}, \
60 columns=userDefined=lengthOfEmbeddingVectors={}'.format(decoder_embedding.shape[0],
61                                                            decoder_embedding.shape[1]))
62 # this returns the embedding vectors of the current input defined by `decoder_input_seq`
63 # (which is actually the output with <GO> at the front)
64 decoder_input_embedded = tf.nn.embedding_lookup(decoder_embedding, decoder_input_seq)
65 print('Shape of embeddings for current input sequence: rows=inputSequenceLength={}, \
66       columns=userDefined=lengthOfEmbeddingVectors={}'\
67       .format(decoder_input_embedded.shape[0], decoder_input_embedded.shape[1]))
68 # add degenerate dimension to emulate batch size
69 decoder_input_embedded = tf.expand_dims(decoder_input_embedded, axis=0)
70 # Output multiRNN
71 decoder_multi_cell = make_multi_cell(RNN_STATE_DIM, RNN_NUM_LAYERS)
72 # The output of the decoder will need to be mapped to a one-hot representation of the vocabulary
73 # this will also be trained
74 # BUT: I'm not sure how the Dense layer output is used since not activation is defined
75 # maybe it's just a linear weighted sum... (?)
76 output_layer = Dense(OUTPUT_NUM_VOCAB, # number of symbols in the output dictionary
77                      kernel_initializer=tf.truncated_normal_initializer(mean=0.0, stddev=0.1))
78 # **** TRAINING DECODER ****
79 # this function manages the input to the decoder for us
80 training_helper = tf.contrib.seq2seq.TrainingHelper(inputs=decoder_input_embedded,
81                                                    sequence_length=(len(padded_symbols_output),), time_major=False)
82 # the training decoder will receive both the state from the encoder `encoder_state`
83 # and the decoder inputs `decoder_input_embedded` via the training helper
84 # this defines the decoder RNN as in the case of `tf.nn.dynamic_rnn`
85 # but this is a bit more complicated because of the input/output relationship that is handled by `training_helper`
86 # and that we need to pass an `encoder_state` to it, as well as the output Dense layer
87 training_decoder = tf.contrib.seq2seq.BasicDecoder(decoder_multi_cell,
88                                                    training_helper, encoder_state, output_layer)
89 # produces the output of the decoder
90 # Specifically it returns [final_outputs, final_state, final_sequence_lengths]
91 training_decoder_output_seq, _, _ = tf.contrib.seq2seq.dynamic_decode(training_decoder,
92                               impute_finished=True, maximum_iterations=len(padded_symbols_output))

```

```

93 sess.run(tf.global_variables_initializer())
94 sess.run(tf.local_variables_initializer())
95 # output of training decoder has shape [shape,numberOfOutputCharacters,sizeOfOutputDictionary]
96 # so for each outputCharacter there is a sequence of numbers of
97 # size=sizeOfOutputDictionary=OUTPUT_NUM_VOCAB that represent each of the possible
98 # characters from the output dictionary
99 print('Shape of training decoder -> [batchSize={}, currentOutputCharacters={}, \
100      sizeOfOutputDictionary={}]' .format(training_decoder_output_seq.rnn_output.eval().shape[0],
101      training_decoder_output_seq.rnn_output.eval().shape[1],
102      training_decoder_output_seq.rnn_output.eval().shape[2]))
103 # **** INFERENCE DECODER ****
104 # Since we don't have an input to the decoder in INFERENCE we will just input the
105 # <GO> token to get the sequence started and start producing the proper output
106 BATCH_SIZE=1
107 start_tokens = tf.tile(tf.constant([output_symbol_to_int['<GO>']],
108      dtype=tf.int32),[BATCH_SIZE],name='start_tokens')
109 # Helper for the inference process. It takes the whole of decoder embeddings
110 # and the integer representation of the <EOS> symbol, which is the 'end_token'
111 # This helper will handle the inputs to the decoder during INFER.
112 inference_helper = tf.contrib.seq2seq.GreedyEmbeddingHelper(embedding=decoder_embedding,
113      start_tokens=start_tokens,end_token=output_symbol_to_int['<EOS>'])
114 # Basic decoder for INFERENCE. It uses the decoder MultiRNN `decoder_multi_cell`
115 # inference_helper takes care of the input (input only the <GO> symbol), the encoder state
116 # which takes the sentence for which we want to have a reply from our chatbot
117 # and the output layer that turns the output into a one-hot representation corresponding to
118 # our vocabulary
119 inference_decoder = tf.contrib.seq2seq.BasicDecoder(decoder_multi_cell,
120      inference_helper,encoder_state,output_layer)
121 # Perform dynamic decoding using the decoder
122 # this is the output, same as in the case of the TRAIN decoder
123 inference_decoder_output_seq, _, _ = tf.contrib.seq2seq.dynamic_decode(inference_decoder,
124      impute_finished=True,maximum_iterations=MAX_CHAR_PER_LINE)
125 # **** TRAINING COST ****
126 # convert output ground truth to tensor
127 outputGroundTruth=tf.convert_to_tensor(padded_symbols_output,dtype=tf.int32)
128 # add batch dummy index
129 batchGroundTruth=tf.expand_dims(outputGroundTruth,axis=0)
130 # mask is used to clip the output sequence to it's maximum allows size
131 # (tf.reduce_max(tf.constant([MAX_CHAR_PER_LINE])))
132 masks = tf.sequence_mask(tf.constant([len(symbols_output)]),MAX_CHAR_PER_LINE,
133      dtype=tf.float32,name='masks')
134 print('Mask for current output: {}'.format(masks.eval()))
135 cost = tf.contrib.seq2seq.sequence_loss(training_decoder_output_seq.rnn_output,
136      batchGroundTruth,masks)
137 print('Cost between model output estimation and grounf truth = {}'.format(cost.eval()))

```

```

*****Input*****
Length of input sequence: 20
*****Embedded Input*****
Shape of embeded input -> rows=lengthInputSequence=20,      columns=userDefined=`ENCODER_EMBEDDING_DIM`=64
Reshaped embedded input -> [batch_size, length_of_sequence,      embedding_dimensions]=[1, 20, 64]
*****Output-States of Encoder*****
Layer 0, hidden state shape (1, 512)
Layer 0, activation state shape (1, 512)
Layer 1, hidden state shape (1, 512)
Layer 1, activation state shape (1, 512)
Unmodified output to decoder (training): [14, 28, 13, 25, 24, 19, 28, 23, 25, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
Modified output for input to decoder (training): [2, 14, 28, 13, 25, 24, 19, 28, 23, 25, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0]
*****Trainable decoder embeddings*****
Shape of decoder embedings are: row=sizeOutputVocabulary=31,      columns=userDefined=lengthOfEmbeddingVectors=64
Shape of embeddings for current input sequence: rows=inputSequenceLength=20,      columns=userDefined=lengthOfEmbeddingVectors=64
Shape of training decoder -> [batchSize=1, currentOutputCharacters=20,      sizeOfOutputDictionary=31]
Mask for current output: [[1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]
Cost between model output estimation and grounf truth = 3.3847382068634033

```

