

# Chapter\_12\_Section\_3\_RankNet\_Utility\_Function

February 9, 2019

## 1 Learn ranking utility function

### 1.1 Practical explanation

This notebook combines the code in Chapter\_12\_Section\_1\_RankNet.ipynb and Chapter\_12\_Section\_2>Loading\_VGG\_Model.ipynb to perform the following: \* Extract first (start\_imgs) and last (end\_imgs) frames of a video demonstrating the folding of clothes. \* Embedd these images using the fully connected output layer of the vgg model shown in Chapter\_12\_Section\_2>Loading\_VGG\_Model.ipynb. \* Use the nn model in Chapter\_12\_Section\_1\_RankNet.ipynb to learn a model to learn a utility function that ranks the last frame (folded) higher than the first frame (unfolded). \* Check that the utility function has learned to score the frames in a video in increasing order as the cloth progress from unfolded to folded.

See “Jointly Learning Grounded Task Structures from Language Instruction and Visual Demonstration” at [http://shukla.io/docs/2016\\_EMNLP.pdf](http://shukla.io/docs/2016_EMNLP.pdf)

### 1.2 Code

Download cloth folding dataset from <https://data.mendeley.com/datasets/c7y3hcrj7z/1>.

Extract the zip. Keep note of where you extract it; we'll call that location DATASET\_DIR.

```
In [1]: import tensorflow as tf
import numpy as np
from vgg16 import vgg16
import glob, os
from scipy.misc import imread, imresize
%matplotlib inline
from matplotlib import pyplot as plt

DATASET_DIR = './data/cloth_folding_rgb_vids'
NUM_VIDS = 45

def get_img_pair(video_id):
    """Extract tuples of start and end images"""
    img_files = sorted(glob.glob(os.path.join(DATASET_DIR, video_id, '*.png')))
    start_img = img_files[0]
    end_img = img_files[-1]
```

```

pair = []
for image_file in [start_img, end_img]:
    img_original = imread(image_file)
    img_resized = imresize(img_original, (224, 224))
    pair.append(img_resized)
return tuple(pair)

start_imgs = []
end_imgs = []
# Extract list of start and end images of unfolded/folded shirts
for vid_id in range(1, NUM_VIDS + 1):
    start_img, end_img = get_img_pair(str(vid_id))
    start_imgs.append(start_img)
    end_imgs.append(end_img)
print('Images of starting state {}'.format(np.shape(start_imgs)))
print('Images of ending state {}'.format(np.shape(end_imgs)))
nExamples=4
fig, ax = plt.subplots(nExamples, 2, figsize=(8, 16))
for i in range(nExamples):
    ax[i, 0].imshow(start_imgs[i])
    ax[i, 0].set_axis_off()
    ax[i, 1].imshow(end_imgs[i])
    ax[i, 1].set_axis_off()
ax[0, 0].set_title('Start')
ax[0, 1].set_title('End')

```

/home/damianos/.conda/envs/tfCPU/lib/python3.5/site-packages/ipykernel/\_\_main\_\_.py:19: DeprecationWarning: `imread` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.

Use ``imageio.imread`` instead.

/home/damianos/.conda/envs/tfCPU/lib/python3.5/site-packages/ipykernel/\_\_main\_\_.py:20: DeprecationWarning: `imresize` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.


Use ``skimage.transform.resize`` instead.

Images of starting state (45, 224, 224, 3)

Images of ending state (45, 224, 224, 3)

Out[1]: Text(0.5, 1.0, 'End')

Start



End



```

In [2]: imgs_plc = tf.placeholder(tf.float32, [None, 224, 224, 3])

In [3]: # nn model to learn to rank two image embeddings
        # See Concept01_ranknet_DC.ipynb

        # this is the size of the embedding
        n_features = 4096
        n_hidden = 10 # n_features * 2

        with tf.name_scope("input"):
            # The input to the ranking nn is the embedding NOT the image
            x1 = tf.placeholder(tf.float32, [None, n_features], name="x1")
            x2 = tf.placeholder(tf.float32, [None, n_features], name="x2")
            dropout_keep_prob = tf.placeholder(tf.float32, name='dropout_prob')

        # Two hidden layers with a relu activation
        # Each hidden layer will take an x1 and x2 inputs
        with tf.name_scope("hidden_layer"):
            with tf.name_scope("weights"):
                w1 = tf.Variable(tf.random_normal([n_features, n_hidden]), name="w1")
                tf.summary.histogram("w1", w1)
                b1 = tf.Variable(tf.random_normal([n_hidden]), name="b1")
                tf.summary.histogram("b1", b1)

            with tf.name_scope("output"):
                h1 = tf.nn.dropout(tf.nn.relu(tf.matmul(x1, w1) + b1), keep_prob=dropout_keep_prob)
                tf.summary.histogram("h1", h1)
                h2 = tf.nn.dropout(tf.nn.relu(tf.matmul(x2, w1) + b1), keep_prob=dropout_keep_prob)
                tf.summary.histogram("h2", h2)

        # Two output layers without an activation function
        # corresponding to each input image embedding
        with tf.name_scope("output_layer"):
            with tf.name_scope("weights"):
                w2 = tf.Variable(tf.random_normal([n_hidden, 1]), name="w2")
                tf.summary.histogram("w2", w2)
                b2 = tf.Variable(tf.random_normal([1]), name="b2")
                tf.summary.histogram("b2", b2)

            with tf.name_scope("output"):
                s1 = tf.matmul(h1, w2) + b2
                s2 = tf.matmul(h2, w2) + b2

```

```

with tf.name_scope("loss"):
    s12 = s1 - s2
    s12_flat = tf.reshape(s12, [-1])

    cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=tf.zeros_like(s12_flat),
                                                             logits=s12_flat + 1)

    loss = tf.reduce_mean(cross_entropy)
    tf.summary.scalar("loss", loss)

with tf.name_scope("train_op"):
    train_op = tf.train.AdamOptimizer(0.001).minimize(loss)

```

WARNING:tensorflow:From <ipython-input-3-6d4668ce4d69>:50: softmax\_cross\_entropy\_with\_logits (Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See @`{tf.nn.softmax_cross_entropy_with_logits_v2}`.

```

In [4]: sess = tf.InteractiveSession()
        sess.run(tf.global_variables_initializer())

In [5]: print('Loading model...')
        vgg = vgg16(imgs_plc, './vggWeights/vgg16_weights.npz', sess)
        print('Done loading!')

```

```

Loading model...
0 conv1_1_W (3, 3, 3, 64)
1 conv1_1_b (64,)
2 conv1_2_W (3, 3, 64, 64)
3 conv1_2_b (64,)
4 conv2_1_W (3, 3, 64, 128)
5 conv2_1_b (128,)
6 conv2_2_W (3, 3, 128, 128)
7 conv2_2_b (128,)
8 conv3_1_W (3, 3, 128, 256)
9 conv3_1_b (256,)
10 conv3_2_W (3, 3, 256, 256)
11 conv3_2_b (256,)
12 conv3_3_W (3, 3, 256, 256)
13 conv3_3_b (256,)
14 conv4_1_W (3, 3, 256, 512)
15 conv4_1_b (512,)
16 conv4_2_W (3, 3, 512, 512)

```

```

17 conv4_2_b (512,)
18 conv4_3_W (3, 3, 512, 512)
19 conv4_3_b (512,)
20 conv5_1_W (3, 3, 512, 512)
21 conv5_1_b (512,)
22 conv5_2_W (3, 3, 512, 512)
23 conv5_2_b (512,)
24 conv5_3_W (3, 3, 512, 512)
25 conv5_3_b (512,)
26 fc6_W (25088, 4096)
27 fc6_b (4096,)
28 fc7_W (4096, 4096)
29 fc7_b (4096,)
30 fc8_W (4096, 1000)
31 fc8_b (1000,)
Done loading!

```

In [6]: *# get embeddings for images*

```

start_imgs_embedded = sess.run(vgg.fc1, feed_dict={vgg.imgs: start_imgs})
end_imgs_embedded = sess.run(vgg.fc1, feed_dict={vgg.imgs: end_imgs})

```

```

idxs = np.random.choice(NUM_VIDS, NUM_VIDS, replace=False)
train_idxes = idxs[0:int(NUM_VIDS * 0.75)]
test_idxes = idxs[int(NUM_VIDS * 0.75):]

```

```

train_start_imgs = start_imgs_embedded[train_idxes]
train_end_imgs = end_imgs_embedded[train_idxes]
test_start_imgs = start_imgs_embedded[test_idxes]
test_end_imgs = end_imgs_embedded[test_idxes]

```

```

print('Train start imgs {}'.format(np.shape(train_start_imgs)))
print('Train end imgs {}'.format(np.shape(train_end_imgs)))
print('Test start imgs {}'.format(np.shape(test_start_imgs)))
print('Test end imgs {}'.format(np.shape(test_end_imgs)))

```

```

Train start imgs (33, 4096)
Train end imgs (33, 4096)
Test start imgs (12, 4096)
Test end imgs (12, 4096)

```

In [7]: *# run nn model that learn to rank start and end images*

```

train_y1 = np.expand_dims(np.zeros(np.shape(train_start_imgs)[0]), axis=1)
train_y2 = np.expand_dims(np.ones(np.shape(train_end_imgs)[0]), axis=1)
for epoch in range(100):
    for i in range(np.shape(train_start_imgs)[0]):
        _, cost_val = sess.run([train_op, loss],

```

```

        feed_dict={x1: train_start_imgs[i:i+1,:],
                    x2: train_end_imgs[i:i+1,:],
                    dropout_keep_prob: 0.5})

    print('{} {}'.format(epoch, cost_val))
    s1_val, s2_val = sess.run([s1, s2], feed_dict={x1: test_start_imgs,
                                                    x2: test_end_imgs,
                                                    dropout_keep_prob: 1})

    print('Accuracy: {}'.format(100 * np.mean(s1_val < s2_val)))

```

```

0. 0.0
Accuracy: 100.0%
1. 0.0
Accuracy: 100.0%
2. 0.0
Accuracy: 100.0%
3. 0.0
Accuracy: 100.0%
4. 0.0
Accuracy: 100.0%
5. 0.0
Accuracy: 100.0%
6. 0.0
Accuracy: 100.0%
7. 0.0
Accuracy: 100.0%
8. 0.0
Accuracy: 100.0%
9. 0.0
Accuracy: 100.0%
10. 0.0
Accuracy: 100.0%
11. 0.0
Accuracy: 100.0%
12. 0.0
Accuracy: 100.0%
13. 0.0
Accuracy: 100.0%
14. 0.0
Accuracy: 100.0%
15. 0.0
Accuracy: 100.0%
16. 0.0
Accuracy: 100.0%
17. 0.0
Accuracy: 100.0%
18. 0.0
Accuracy: 100.0%
19. 0.0

```

Accuracy: 100.0%  
20. 0.0  
Accuracy: 100.0%  
21. 0.0  
Accuracy: 100.0%  
22. 0.0  
Accuracy: 100.0%  
23. 0.0  
Accuracy: 100.0%  
24. 0.0  
Accuracy: 100.0%  
25. 0.0  
Accuracy: 100.0%  
26. 0.0  
Accuracy: 100.0%  
27. 0.0  
Accuracy: 100.0%  
28. 0.0  
Accuracy: 100.0%  
29. 0.0  
Accuracy: 100.0%  
30. 0.0  
Accuracy: 100.0%  
31. 0.0  
Accuracy: 100.0%  
32. 0.0  
Accuracy: 100.0%  
33. 0.0  
Accuracy: 100.0%  
34. 0.0  
Accuracy: 100.0%  
35. 0.0  
Accuracy: 100.0%  
36. 0.0  
Accuracy: 100.0%  
37. 0.0  
Accuracy: 100.0%  
38. 0.0  
Accuracy: 100.0%  
39. 0.0  
Accuracy: 100.0%  
40. 0.0  
Accuracy: 100.0%  
41. 0.0  
Accuracy: 100.0%  
42. 0.0  
Accuracy: 100.0%  
43. 0.0



Accuracy: 100.0%  
44. 0.0  
Accuracy: 100.0%  
45. 0.0  
Accuracy: 100.0%  
46. 0.0  
Accuracy: 100.0%  
47. 0.0  
Accuracy: 100.0%  
48. 0.0  
Accuracy: 100.0%  
49. 0.0  
Accuracy: 100.0%  
50. 0.0  
Accuracy: 100.0%  
51. 0.0  
Accuracy: 100.0%  
52. 0.0  
Accuracy: 100.0%  
53. 0.0  
Accuracy: 100.0%  
54. 0.0  
Accuracy: 100.0%  
55. 0.0  
Accuracy: 100.0%  
56. 0.0  
Accuracy: 100.0%  
57. 0.0  
Accuracy: 100.0%  
58. 0.0  
Accuracy: 100.0%  
59. 0.0  
Accuracy: 100.0%  
60. 0.0  
Accuracy: 100.0%  
61. 0.0  
Accuracy: 100.0%  
62. 0.0  
Accuracy: 100.0%  
63. 0.0  
Accuracy: 100.0%  
64. 0.0  
Accuracy: 100.0%  
65. 0.0  
Accuracy: 100.0%  
66. 0.0  
Accuracy: 100.0%  
67. 0.0

Accuracy: 100.0%  
68. 0.0  
Accuracy: 100.0%  
69. 0.0  
Accuracy: 100.0%  
70. 0.0  
Accuracy: 100.0%  
71. 0.0  
Accuracy: 100.0%  
72. 0.0  
Accuracy: 100.0%  
73. 0.0  
Accuracy: 100.0%  
74. 0.0  
Accuracy: 100.0%  
75. 0.0  
Accuracy: 100.0%  
76. 0.0  
Accuracy: 100.0%  
77. 0.0  
Accuracy: 100.0%  
78. 0.0  
Accuracy: 100.0%  
79. 0.0  
Accuracy: 100.0%  
80. 0.0  
Accuracy: 100.0%  
81. 0.0  
Accuracy: 100.0%  
82. 0.0  
Accuracy: 100.0%  
83. 0.0  
Accuracy: 100.0%  
84. 0.0  
Accuracy: 100.0%  
85. 0.0  
Accuracy: 100.0%  
86. 0.0  
Accuracy: 100.0%  
87. 0.0  
Accuracy: 100.0%  
88. 0.0  
Accuracy: 100.0%  
89. 0.0  
Accuracy: 100.0%  
90. 0.0  
Accuracy: 100.0%  
91. 0.0

```

Accuracy: 100.0%
92. 0.0
Accuracy: 100.0%
93. 0.0
Accuracy: 100.0%
94. 0.0
Accuracy: 100.0%
95. 0.0
Accuracy: 100.0%
96. 0.0
Accuracy: 100.0%
97. 0.0
Accuracy: 100.0%
98. 0.0
Accuracy: 100.0%
99. 0.0
Accuracy: 100.0%

```

```

In [8]: # See how well our nn model has learned to rank sequential images of a video
        # NOT just start and end frames as it was trained on
        def get_img_seq(video_id):
            img_files = sorted(glob.glob(os.path.join(DATASET_DIR, video_id, '*.png')))
            imgs = []
            for image_file in img_files:
                img_original = imread(image_file)
                img_resized = imresize(img_original, (224, 224))
                imgs.append(img_resized)
            return imgs

        imgs = get_img_seq('2')
        print(np.shape(imgs))

(37, 224, 224, 3)

```

```

/home/damianos/.conda/envs/tfCPU/lib/python3.5/site-packages/ipykernel/__main__.py:7: Deprecat
`imread` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.
Use ``imageio.imread`` instead.
/home/damianos/.conda/envs/tfCPU/lib/python3.5/site-packages/ipykernel/__main__.py:8: Deprecat
`imresize` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.
Use ``skimage.transform.resize`` instead.

```

```

In [9]: # Get embeddings
        imgs_embedded = sess.run(vgg.fc1, feed_dict={vgg.imgs: imgs})
        scores = sess.run([s1], feed_dict={x1: imgs_embedded, dropout_keep_prob: 1})

In [10]: plt.figure()
          plt.title('Utility of cloth-folding over time')

```

```
plt.xlabel('time (video frame #)')  
plt.ylabel('Utility')  
plt.plot(scores[-1])
```

Out[10]: [<matplotlib.lines.Line2D at 0x7f25b4371cc0>]



sess.close()